Web Exercise 04: Working with Amazon AWS Educate Starter

Due Date: **October 18** (Tuesday), 2016. 3:30pm on Blackboard.

Grade: **15 points.**

1. Create Your AWS Free Account and an EC2 Instance.

Go to the following URL to establish your Amazon AWS Educate Starter account:

https://aws.amazon.com/education/awseducate/apply/

We will be creating an Amazon Web Services (**AWS**) account associated with San Diego State University, so click the "**Students: [Apply for AWS Educate for Students]**" button (the right one):

Then, select your role as "Student". Click on [Next].

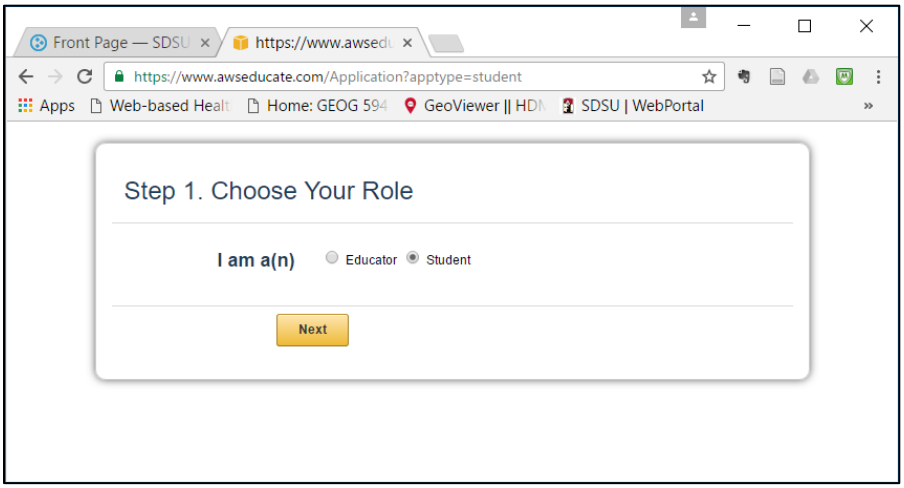

Fill out your application. Type "**San Diego State University**" in the Institution Name. Make sure to use your SDSU mail addresses (such as @rohan.sdsu.edu, @mail.sdsu.edu, or @sdsu.edu) and select the "**Click here to select an AWS Educate Starter Account**." (Ignore the Promotion Code).

Note: **DO NOT SELECT the first option** (Enter an AWS Account ID) since it will require a credit card number to register an AWS account).

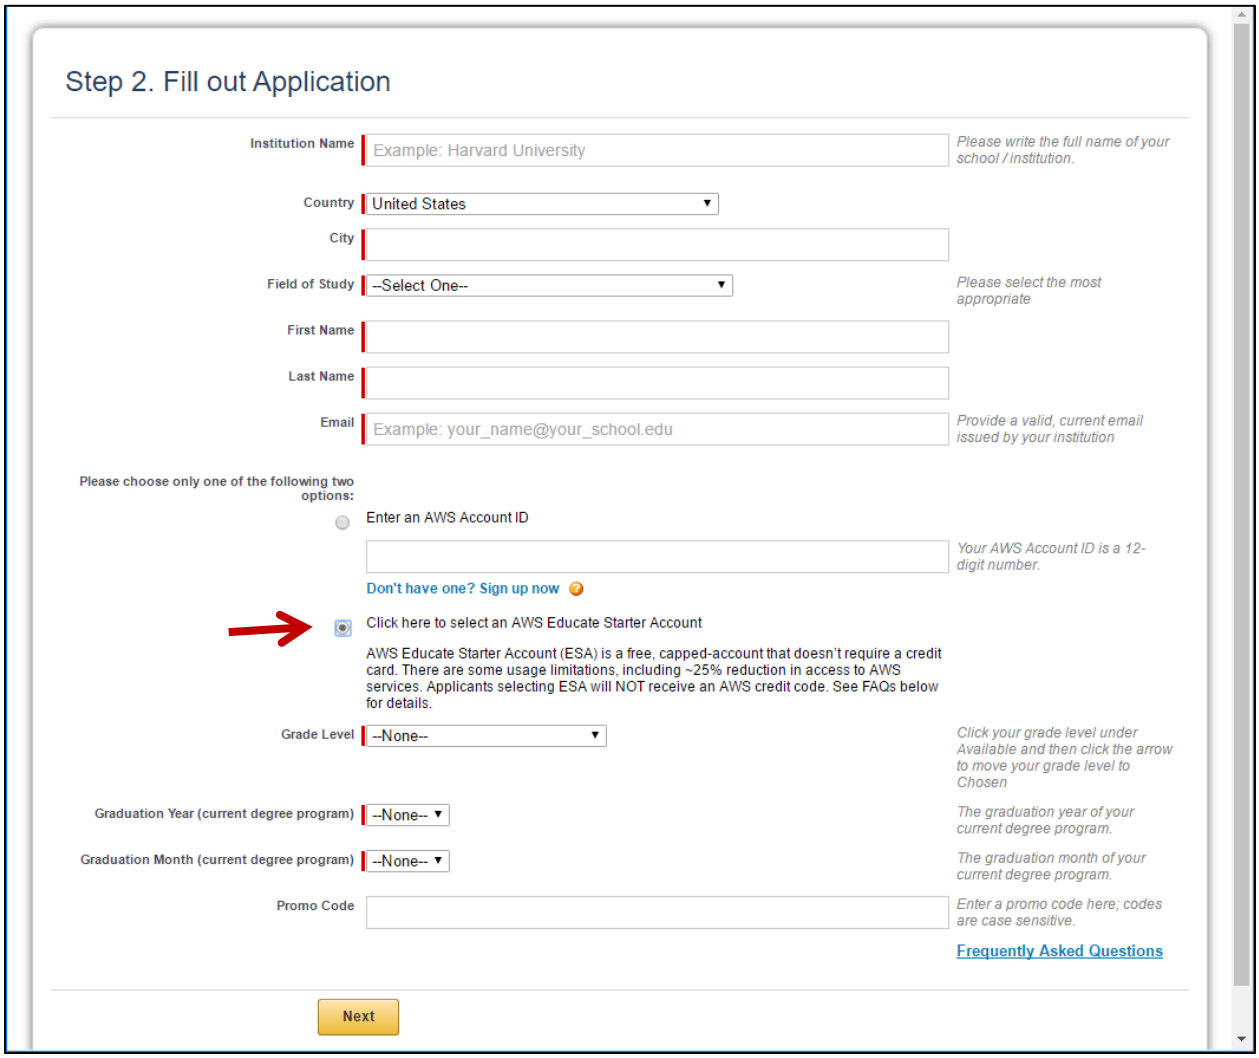

After you submit the request, you will receive a confirmation email from Amazon first. After one or two days, your application will be approved by Amazon and you will be able to set up your user name and password to access your free AWS Educate Starter account.

Once you have received an email confirming that your account has been approved, log into the AWS Educate Portal by going to the following URL:

https://www.awseducate.com/microsite/SiteLogin

Once you log in, you should see the follow page. Click the "Go to your AWS Educate Starter Account" button.

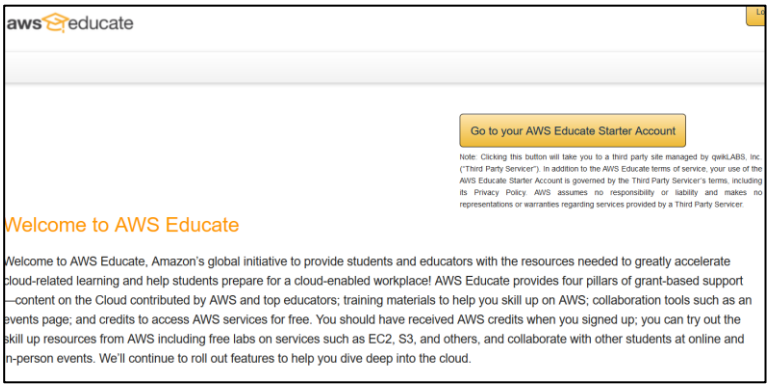

The Qwik Labs Dashboard allows you to view your Amazon accounts and view tutorials related to Amazon's services. Since you signed up for an account that is associated with San Diego State University, you should see one account. Click the links that says "**AWS Educate Stater Account 75.**

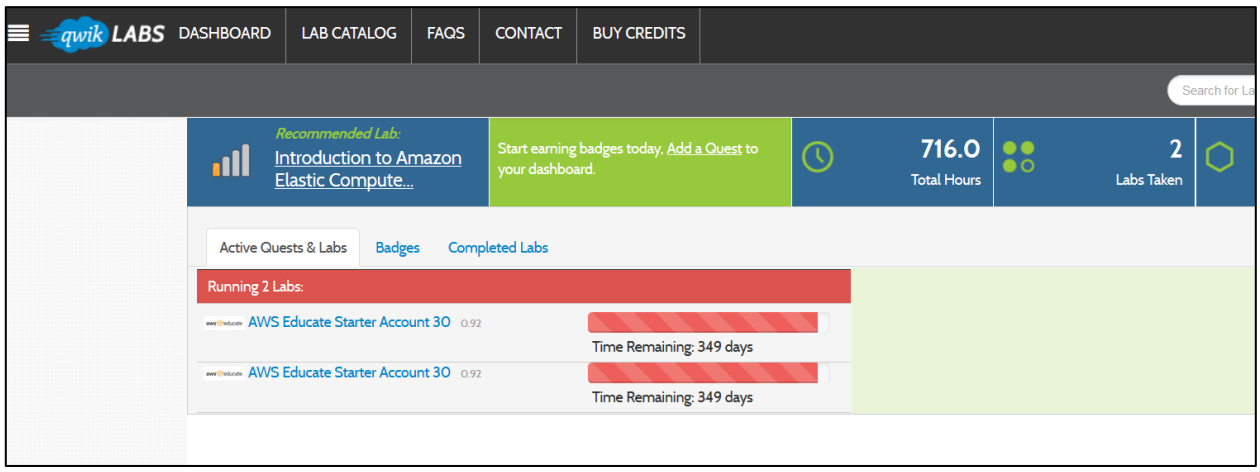

The next step is to **start** the lab. After you start the lab, the "**Open Console**" button will become active. Click on the "**Open Console**" to allow us to configure our Amazon EC2 instances.

This will take us to the Amazon Web Services dashboard. Form here you can configure and administer your Amazon services. For this lab, we are only concerned with Amazon's EC2 service.

Click the "**EC2**" icon under the "Compute" menu.

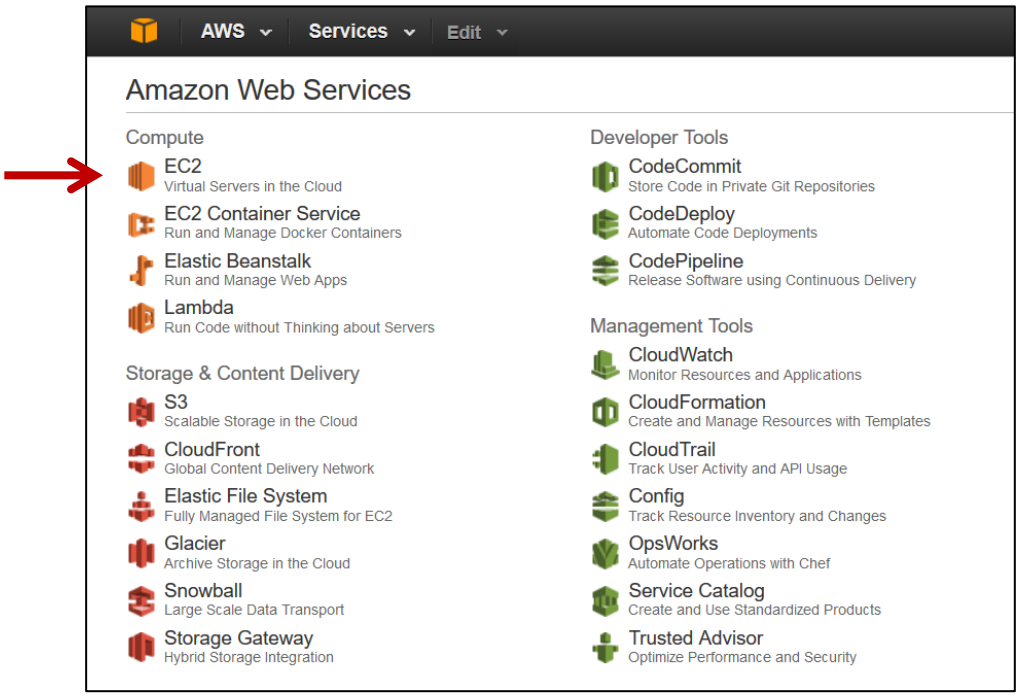

This dashboard shows us any instances of EC2 that we have created. Since we are opening a new account, there should be none present. We want to create a new instance by clicking the "**Launch Instance**" button.

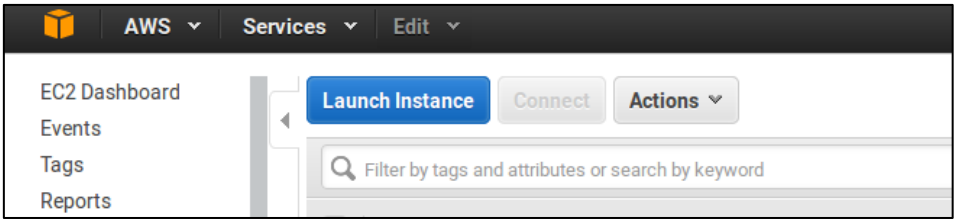

On the next page, we can select the operating system for our instance. For this lab, we want to create an instance of "**Microsoft Windows Server 2012 R2 Base**", so click the "Select" button for this option.

On "Step 2" we can choose the size of our instance. This allows us to configure the computing power and other system resources for our instance. Since we are using a free account, we must select the "**t2 micro" option (free tier eligible)**. Then press "**Review and Launch**" button.

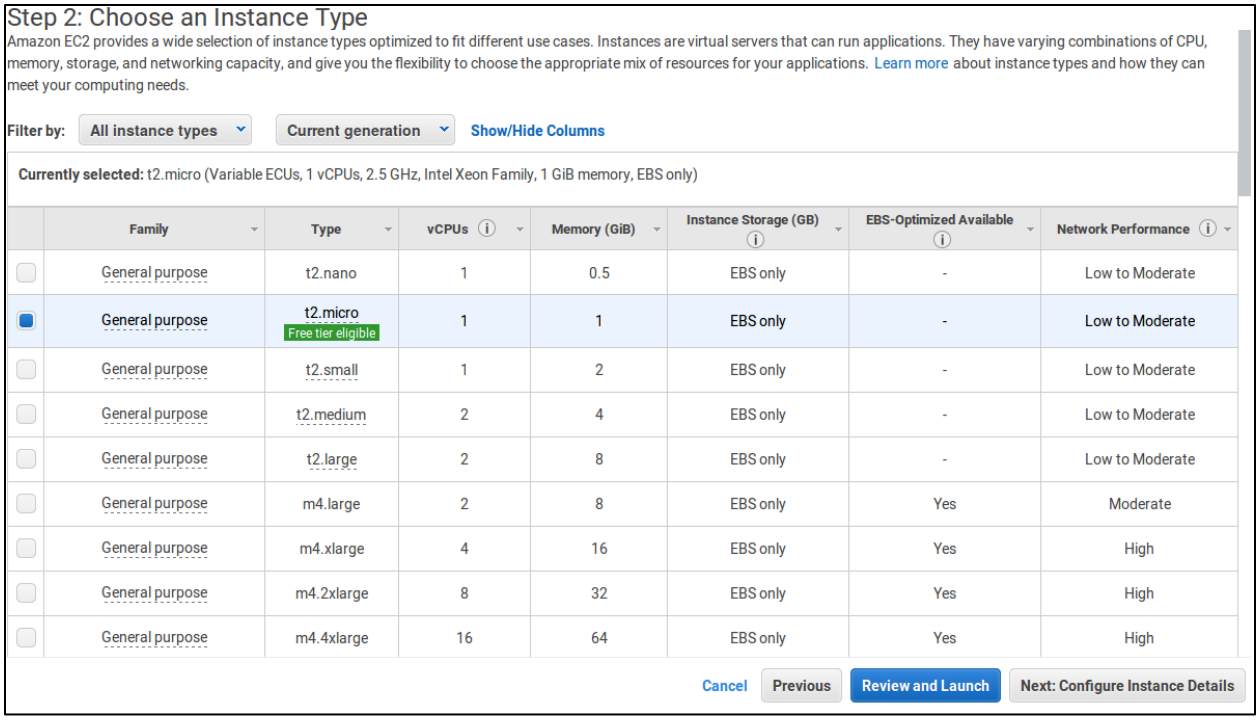

Step 2: Choose an Instance Type

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases. Instances are virtual servers that can run applications. They have varying combinations of CPU, memory, storage, and networking capacity, and give you the flexibility to choose the appropriate mix of resources for your applications. Learn more about instance types and how they can meet your computing needs.

Filter by: All instance types | Current generation | Show/Hide Columns

Currently selected: t2.micro (Variable ECUs, 1 vCPUs, 2.5 GHz, Intel Xeon Family, 1 GiB memory, EBS only)

	Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
☐	General purpose	t2.nano	1	0.5	EBS only	-	Low to Moderate
☑	General purpose	t2.micro Free tier eligible	1	1	EBS only	-	Low to Moderate
☐	General purpose	t2.small	1	2	EBS only	-	Low to Moderate
☐	General purpose	t2.medium	2	4	EBS only	-	Low to Moderate
☐	General purpose	t2.large	2	8	EBS only	-	Low to Moderate
☐	General purpose	m4.large	2	8	EBS only	Yes	Moderate
☐	General purpose	m4.xlarge	4	16	EBS only	Yes	High
☐	General purpose	m4.2xlarge	8	32	EBS only	Yes	High
☐	General purpose	m4.4xlarge	16	64	EBS only	Yes	High

Cancel | Previous | Review and Launch | Next: Configure Instance Details

Once we have launched our instance, a page will tell us our launch status. In order to view our instances, click the hyperlink containing the instance number (here it's i-003becc2f4cb5cfde).

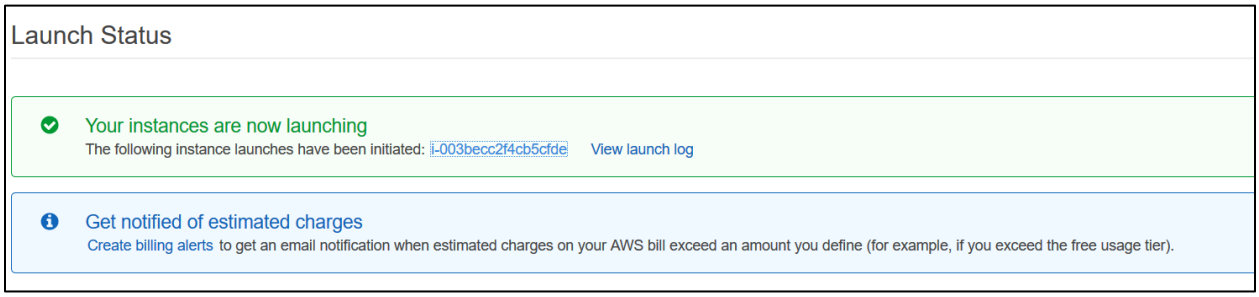

Launch Status

Your instances are now launching

The following instance launches have been initiated: i-003becc2f4cb5cfde View launch log

Get notified of estimated charges

Create billing alerts to get an email notification when estimated charges on your AWS bill exceed an amount you define (for example, if you exceed the free usage tier).

Now that we have created an instance, we need to configure it in order to make a connection through remote desktop. Right click the instance we just created and choose the "**Get Windows Password**" option.

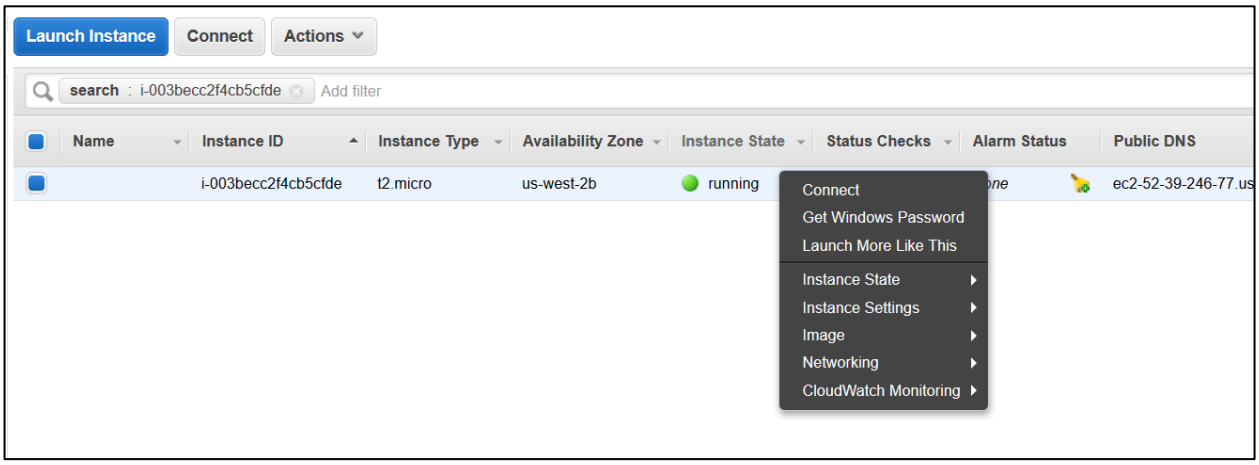

A popup window will be shown that allows you to create a key that will be used to retrieve your windows password. Select the "Create a new key pair" option and give the key a name (in this example, we have named it "**Yourname-Key**" (for example, "Ming-Key"). Next, click the "Download Key Pair" button and save the file to your local computer.

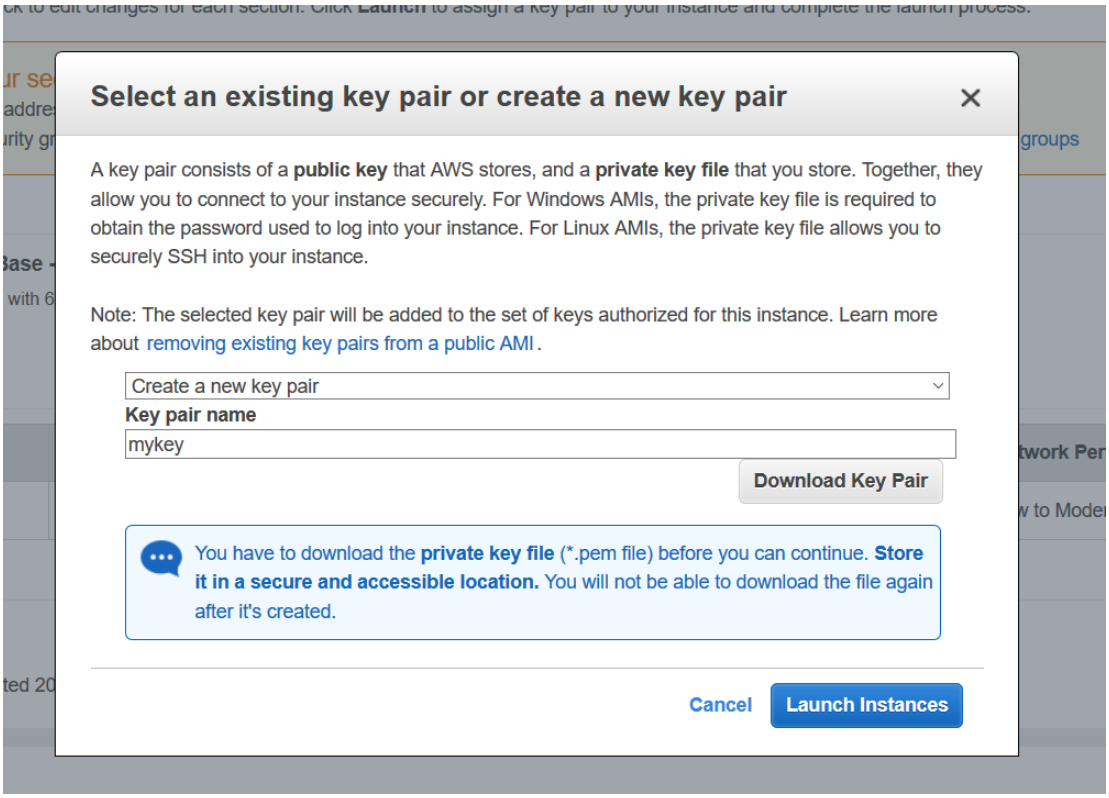

As before, we can right click on our instance and select the "Get Windows Password" option. This will allow us to upload the file we downloaded.

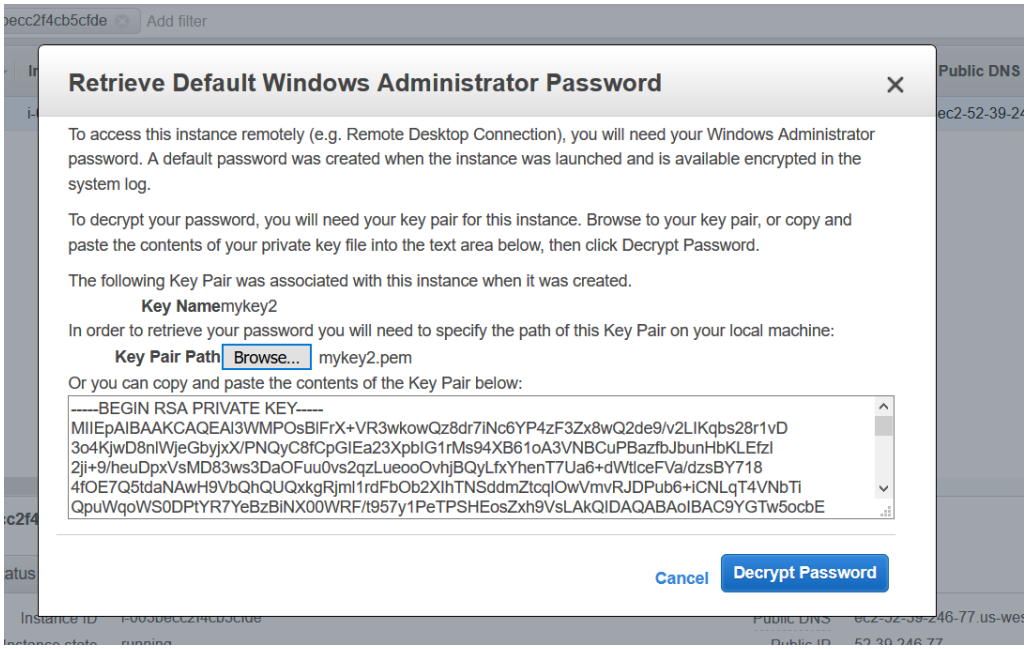

Click the "**Decrypt Password**" button, which will give us the server name, user name and password for our EC2 instance. These will be used for logging into our instance through remote desktop.

Open a Remote Desktop Connection from your local computer. Cut and paste the DNS server information for your instance. When you click "Connect", you will be prompted for a user name and password. Enter the information (password) from the previous step to log into the Windows Server.

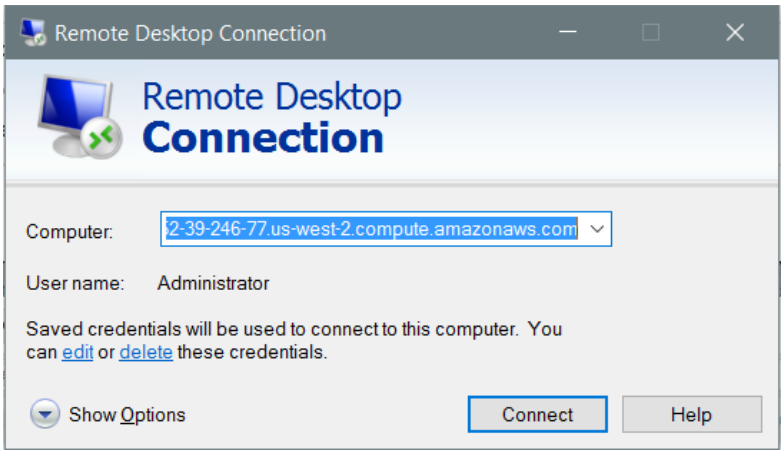

Another option to use the Remote Desktop Login is to click on the [Connect] button on the top of the window. Then click on [Download Remote Desktop File]. After downloading, double click on the file to open the Remote Desktop.

After the login, you can see the Remote Desktop as the following:

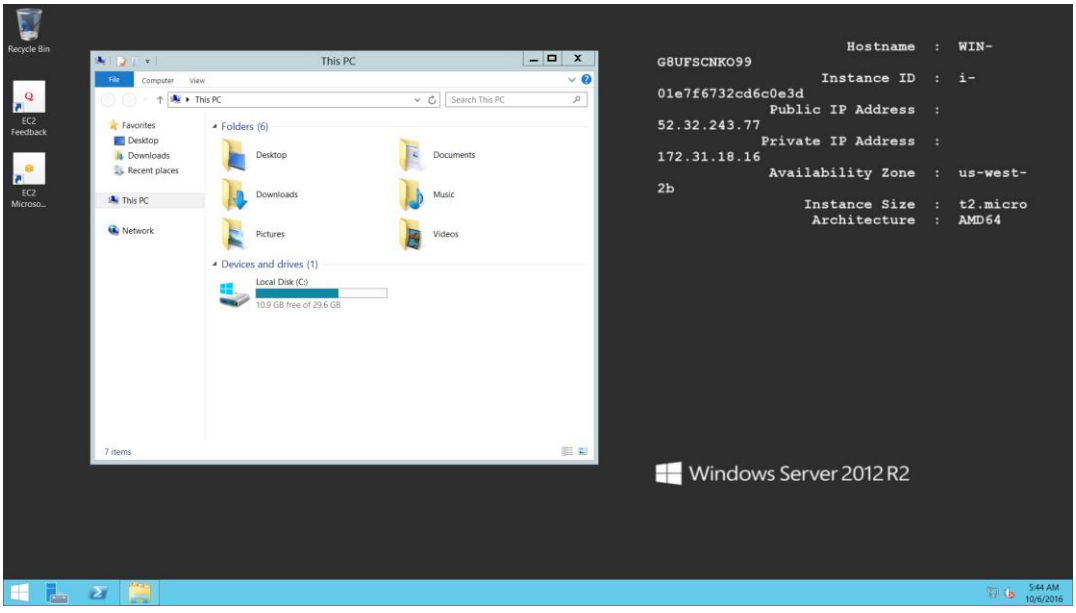

The next step to install software and a NoSQL databases that will allow us to collect tweet and store them in a database.

You can "**Disconnect**" your EC2 using the Remote Desktop (by clicking on the "X" on the top bar). The Server will keep running when the desktop is disconnected. You can just reconnect it when you need to continue your tasks next time.

2. Install MongoDB (NoSQL) Database

After using the Remote Desktop Login to your EC2 instance. Open the **Internet Explorer** inside your EC2 instance.

First we want to install MongoDB. This is a database that is optimized for large collections of data that are often unstructured (for more information: https://en.wikipedia.org/wiki/MongoDB). To download MongoDB, go to https://www.mongodb.com/download-center?jmp=nav#community (from IE) and download the Windows installer for x64. You may receive a Security Alert indicating that "You are about to view pages…" You can check the box and then click on [OK]. When you access the MongoDB website, you can add these URL into the "Trusted sites" and then continue the download process.

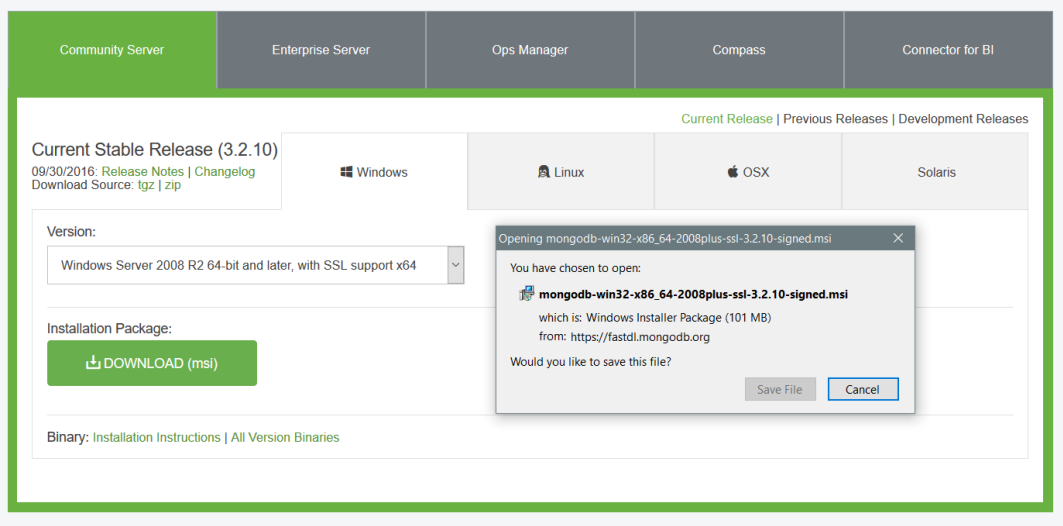

Download and Run the installer and when you get to the "Choose Setup Type" page, select the "**Custom**" option.

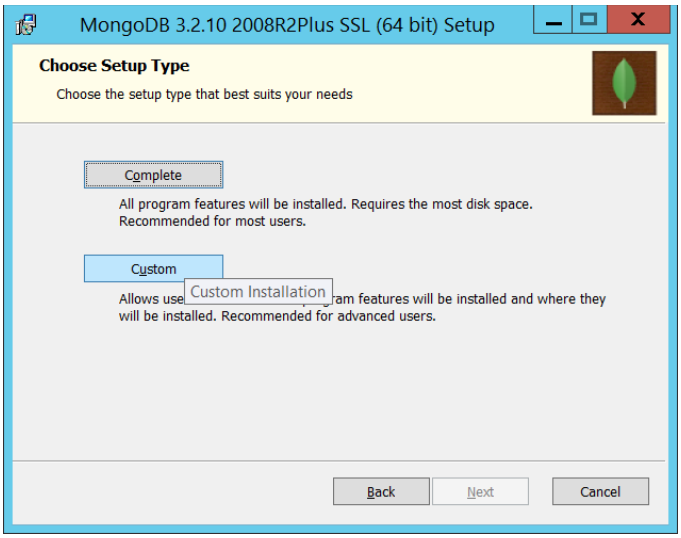

In the "Custom Setup" page, change the Location by clicking on "Browse.." and type in the Folder name "**C:\MongoDB\Server\3.2**", then click on **OK**. This will simplify the installation and running of MongoDB. Continue with the setup selecting the default options. (click on NEXT and then). Click on the [Finish] button to Exit.

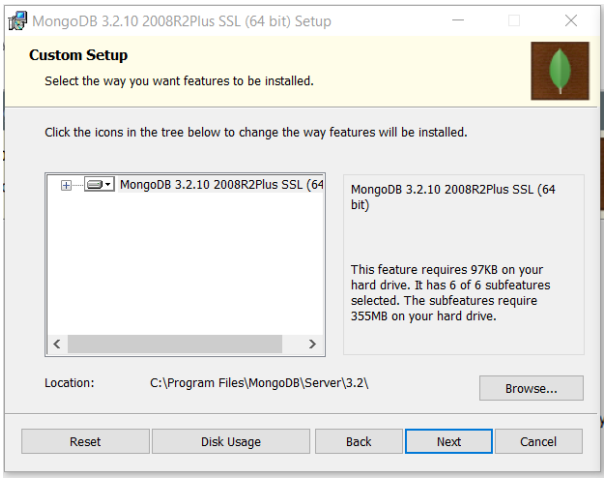

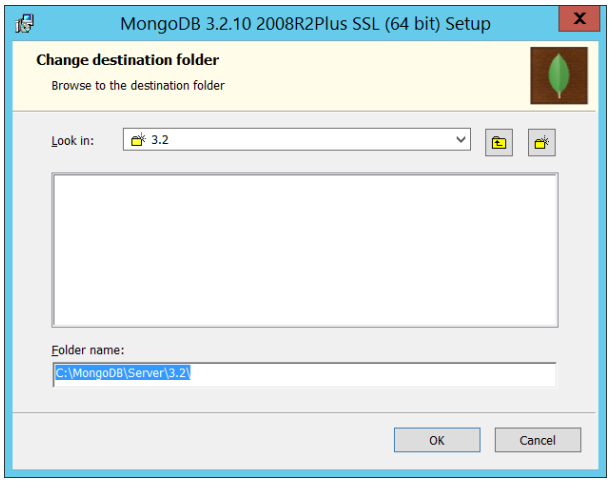

Open the File Explorer (Right Click on the Start-up Menu → select "File Explorer".

Create a new folder inside the **C:** drive called "**data**." Within that folder, create a sub folder called "**db**". This is the folder in which data will be stored by MongoDB. Please make sure that the "**C:/data/db**" is at the same level of "MongoDB" folder. Otherwise, the MongoDB may not be able to generate a database.

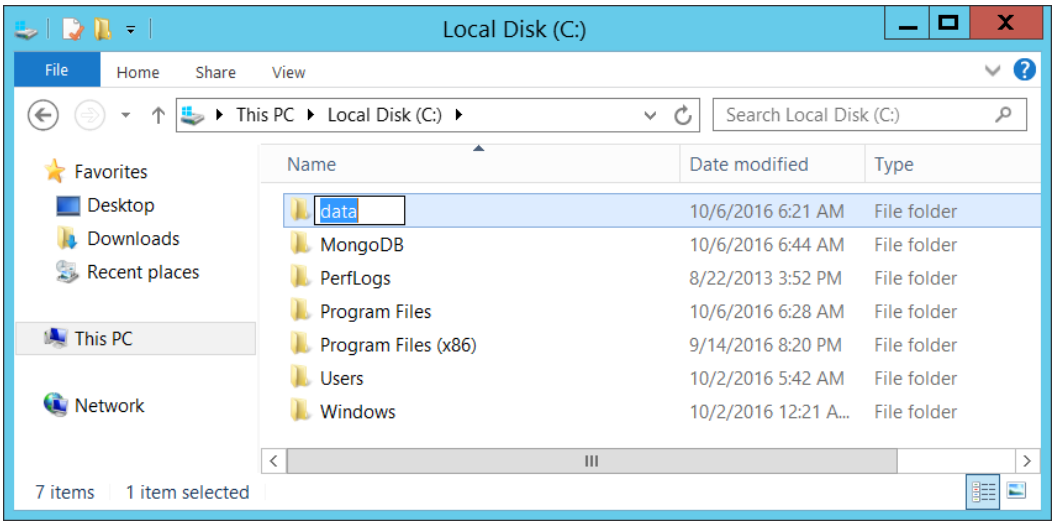

Start MongoDB by opening the [**command line interface] (right click on Start-up menu)** or the Windows **PowerShell** (3) (blue icon on the Task Bar) and navigating to the folder in which you installed MongoDB (Type "**cd C:\MongoDB\server\3.2**". Navigate to the "bin" folder (Type "**cd bin**"). Then, type the command "**.\mongod**" to start the MongoDB. **Leave this window open, so MongoDB can continue to work on your EC2 machine. DO NOT close this window.**

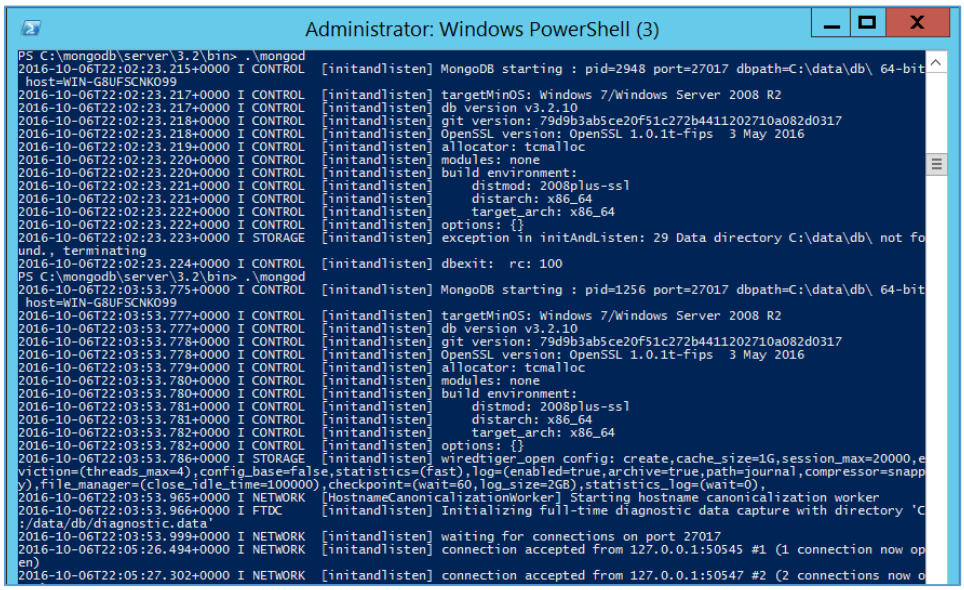

3. Install RoboMongo (a Graphic User Interface for MongoDB).

It will be helpful to install a graphical user interface for administering our MongoDB databases. Go the following page and select the 64-bit .ngo to install this tool into your EC2. (Use all default setting).

https://robomongo.org/download

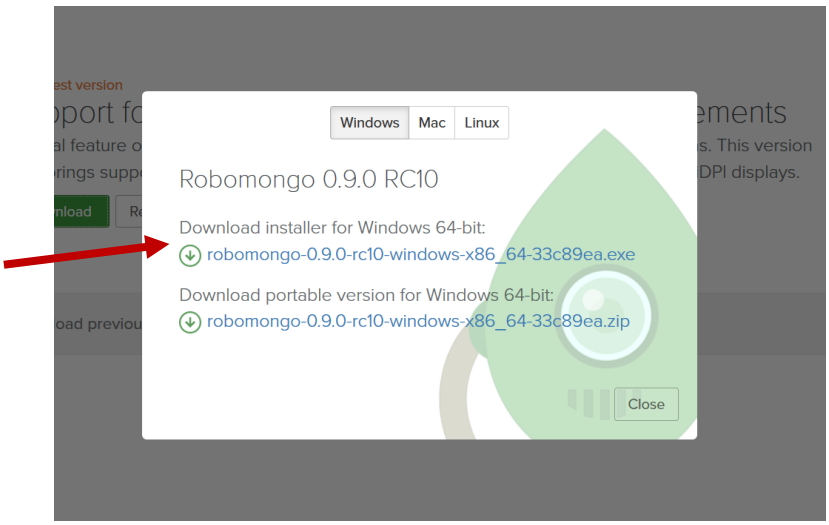

4. Install Python and related Libraries.

Now we want to install Python and libraries that will allow us to download tweets and insert them into MonoDB.

First download **Anaconda**, which is a distribution of Python that includes many libraries that are useful for data science and scientific computing. Go to the following page to download the installer for Anaconda and click the button for **Python 2.7, 64-Bit Installer**. (Use all default setting).

https://www.continuum.io/downloads

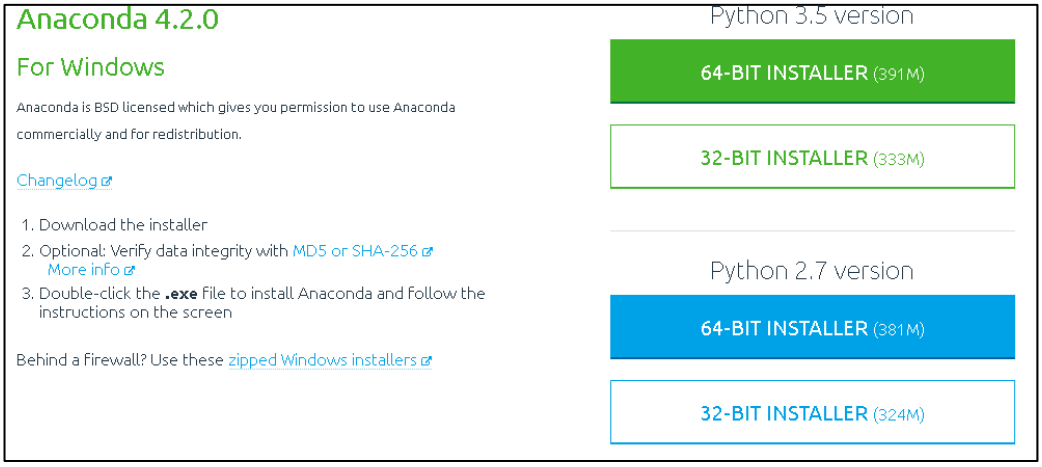

Next, open the Command Line or the PowerShell and enter the following command to install "tweepy," which is a python library that will allow us to collect tweets from the Twitter API:

pip install tweepy

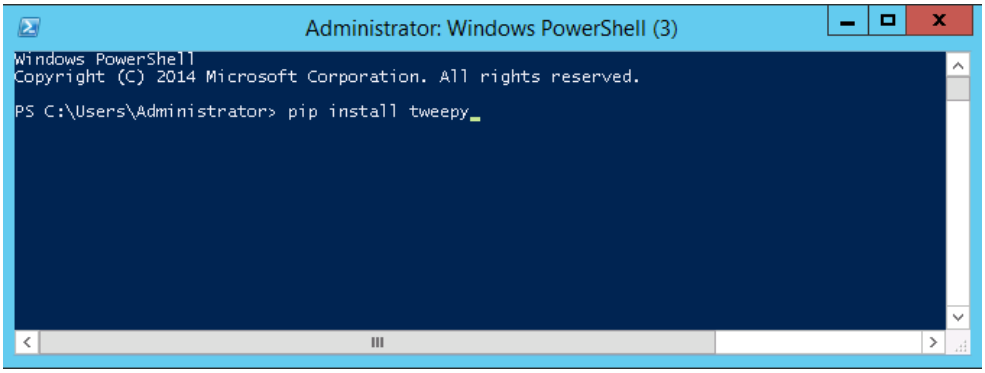

Next, we want to install “pymongo,” which is a Python library that allows us to interface with MongoDB. Type the following command to install “pymongo”:

pip install pymongo

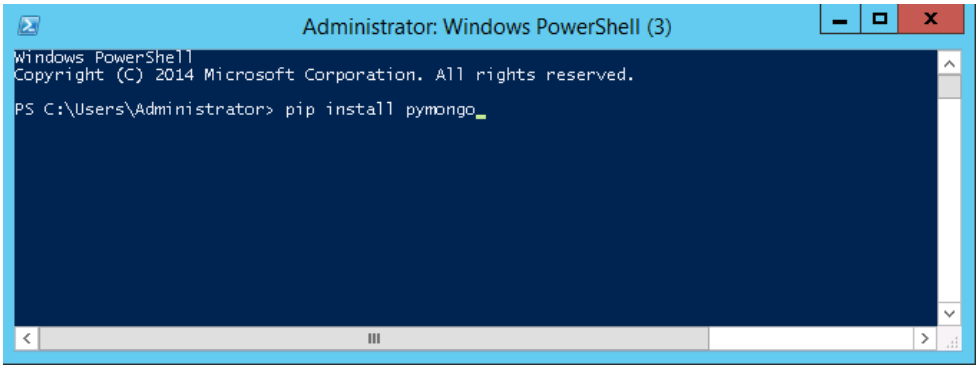

Next we will create a Python script that downloads tweets from the Twitter API and places them into MongoDB.

5. Get the Twitter API Keys/Tokens.

First, we will need to sign up for a Twitter account (if you don’t already have one) and obtain some information that will allow us to connect to the Twitter API. Make sure to add your mobile phone number to your Twitter Profile before the following step.

Here are the simple procedure to get your Twitter API key/Tokens.

1. Log on to Twitter Developer (**https://dev.twitter.com/**) and sign in with your Twitter account.
2. Next, use the same web browser to open this URL**:** **https://apps.twitter.com/**
3. Click on “**Create New App**” button.
4. Fill out the application form. (for the Website URL, If you don't have a URL yet, just put a placeholder here but remember to change it later.) Make sure to check the YES for the Developer Agreement.

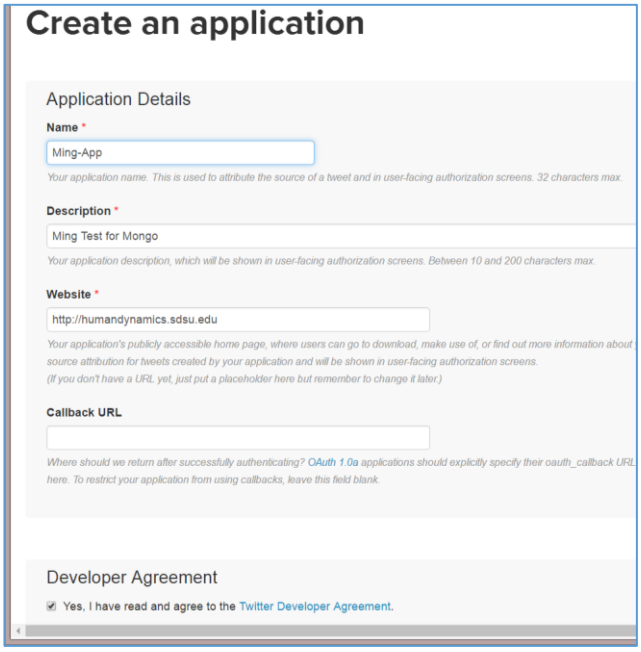

5. Click on "**Create your Twitter application**" button. (Note: you need to add your mobile phone to your Twitter profile before creating an application.)

After that, you will see the following page:

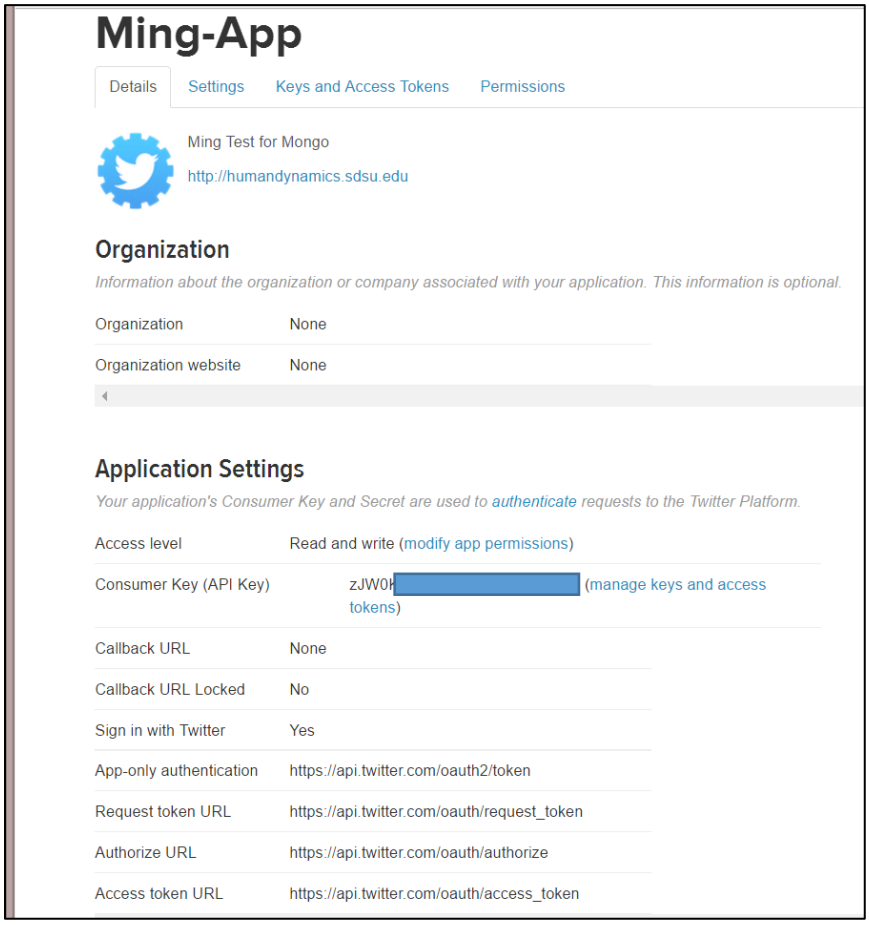

Click on "**Keys and Access Tokens**" Tab on the upper menu (under your application name). Save the Consumer key (API Key) and Consumer secret (API Secret) into your local file (in Word or text file). In the bottom of the page, click on "**Create my access token**". Copy and Save your Access Token and Access Token Secret into your local file. You will need to use the four items in the next step.

6. Create a Python Program to Collect Tweets.

Next, we want to create a Python script, called "**get_tweets.py**", which establishes a connection to the Twitter API, downloads tweets, and stores them in MongoDB:

Open a programing text editing tool, like **notepad++ (https://notepad-plus-plus.org/)** or **ATOM** (https://atom.io/).

Create the following program and save it (as "**get_tweets-yourname.py**") into your EC2 instance C:\Scripts folder (You will need to create this new folder first). Make sure to type the Twitter **API KEY, Secret, Token, and Token_Secret** into this python program.

```
import tweepy
import pymongo

# create a connection to MongoDB
client = pymongo.MongoClient()
# create a database called 'data'
database = client['data']
# create a collection (ie, a table) called 'tweets'
collection = database['tweets']


# The consumer keys can be found on your application's Details
# page located at https://dev.twitter.com/apps (under "OAuth settings" and "Your access token")
consumer_key = '<YOUR_CONSUMER_KEY>'
consumer_secret = '<YOUR_CONSUMER_SECRET>'
access_token = '<YOUR_ACCESS_TOKEN>'
access_token_secret = '<YOUR_ACCESS_TOKEN_SECRET>'

# authenticate your key with oauth
auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.secure = True
auth.set_access_token(access_token, access_token_secret)

# setup the API with your Keys
# also, format the return in JSON
api = tweepy.API(auth, parser=tweepy.parsers.JSONParser())

# Calling the search method, with several parameter
# q: querying keyword
# count: limits the numbers of return
result = api.search(q='penguin', count=100)

# grab the tweets from the search result
tweets = result['statuses']

# insert these tweets into the 'tweets' collection
collection.insert(tweets)


# now, let's query the tweets we just stored in the collection
# collection.find() willl return all records in the collection
records = collection.find()

for record in records:
    # get user and tweet text (note: text must be converted to ascii to display in the console)
    user = record['user']['screen_name']
    tweet = record['text'].encode('ascii', 'ignore')

    print 'User:', user
    print 'Tweet:', tweet
    print '\n'
```

This program will search 100 Twitter messages containing "penguin" and then save the results into the MongoDB database. **Please choose our own preferred search keyword (replacing "penguin" with your own keyword, such as "SDSU" or "iPhone".**

To run this script, open a command prompt and navigate to the folder where you saved your script into a new folder "**C:\Scripts**" (you will create this new folder first). Assuming you have called your file "**get_tweets-yourname.py**", you can run the following command to download 100 tweets related to your keyword and store them in MongoDB:

python C:\Scripts\get_tweets-yourname.py

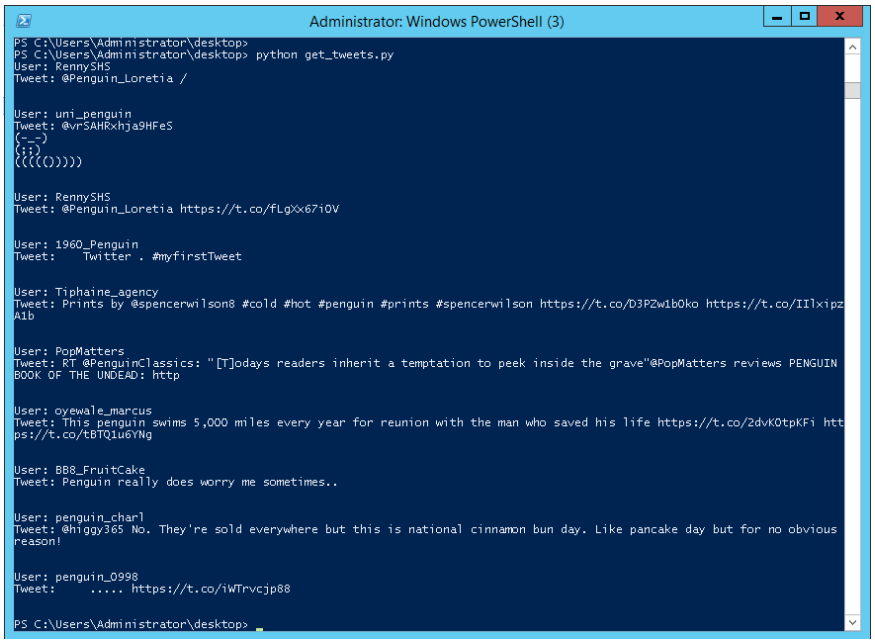

Note: Some of the tweets in the screenshot above do not have the selected keyword in the text. This is because the Twitter API will return tweets if the user's screen name also contains the keyword (even if the text does not).

Next, we can open up RoboMongo. The first time this program is opened it will ask you to create a MongoDB connection. Click the "**Create**" link and enter a name for our connection. The default settings (localhost) should be fine to configure this connection. Click "**Save**" and then connect to your connection.

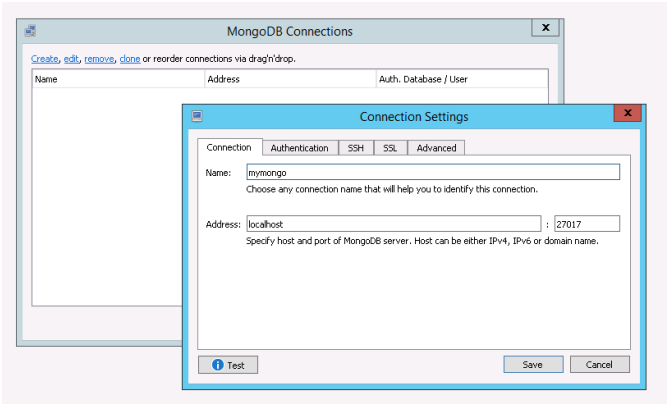

In the Python script we created above, we stored the tweets in the "tweets" collection within the "data" database. Navigate to this collection within RoboMongo. From here we can inspect the tweets we stored using the Python script: System → data → Collection(1) → tweets

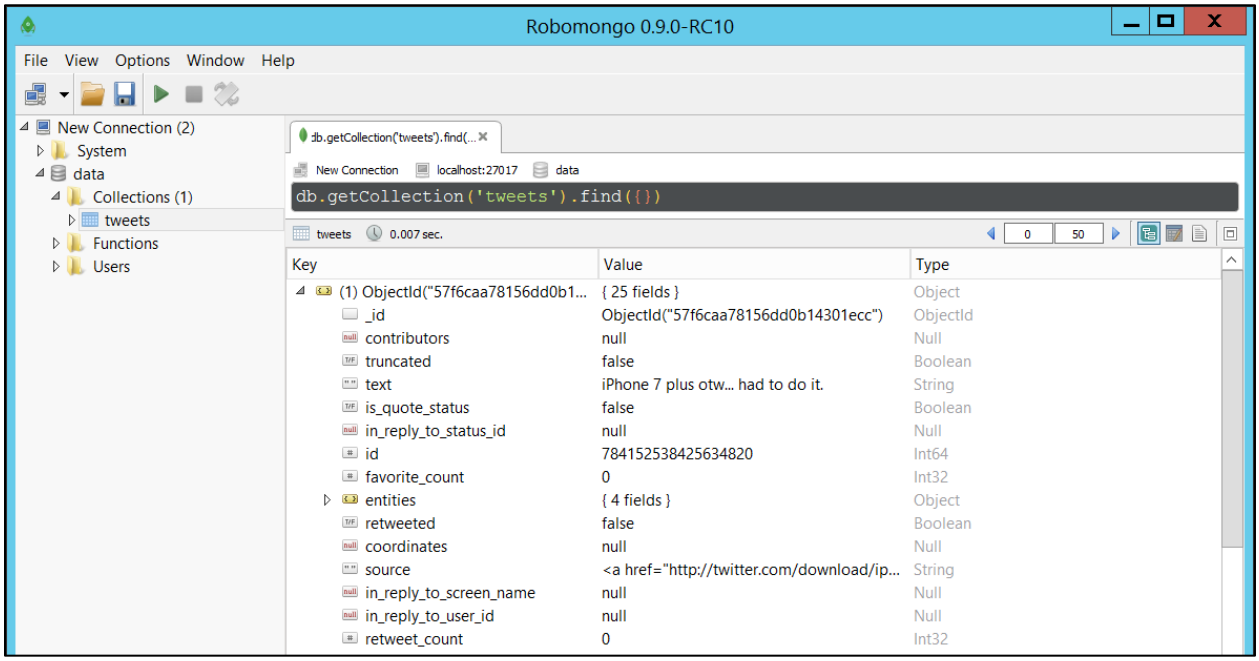

Take a close look at the content of one selected tweet inside the MongoDB. What are the KEYS in the databases? What are the Values associated with each KEY? What type of Value are they? Read the Data Model Introduction in MongoDB online document:

https://docs.mongodb.com/manual/core/data-modeling-introduction/

After finishing all steps and tasks, you can just keep your free EC2 Instance running, since it is a free-tier instance. If you are paying the usage of EC2 instance, you can "STOP" it when you are not running the server to save money (similar to "shutdown" your server). When you need to re-use the Instance, you can "Start" the instance again.

If you select "STOP" for your instance, you will see a warning message that "Any data on the ephemeral storage of your instance will be lost". Since we are not using any ephemeral storage in this exercise. It is fine to STOP the instance. (Note, but you will need to "re-launch" the MongoDB database when you re-START this instance next time).

After finishing this Web Exercise, Please use your own words to answer the following questions (next page): (DO NOT COPY any web resources or Wikipedia texts. We will check your answers with Blackboard tools to verify that your responses are uniquely yours.) By submitting your answers (paper) to Blackboard, you agree: (1) that you are submitting your paper to be used and stored as part of the SafeAssign™ services in accordance with the Blackboard Privacy Policy; (2) that your institution may use your paper in accordance with your institution's policies; and (3) that your use of SafeAssign will be without recourse against Blackboard Inc. and its affiliates.

SafeAssign accepts files in .doc, .docx, .docm, .ppt, .pptx, .odt, .txt, .rtf, .pdf, and .html file formats only. Files of any other format will not be checked through SafeAssign.

LAB-4 Additional Assignment:

1. What are the advantages and disadvantage of using Amazon EC2 for building a research project website?
2. What are the differences between traditional SQL databases and NoSQL databases?
3. What is your selected keyword in the **get_tweets-yourname.py** to collect tweets? How many tweets did you get? **List the content of first 3 tweets in your report. Discuss the search results and how to improve the keyword search for your subject.**

Please submit your LAB-4 Answers (in a MS Word or a PDF file format only) to the Blackboard System BEFORE the DUE DATE/TIME.